

Terapias Bogota

AI-Assisted Lead Triage

Interaction • Evaluation • Refinement • Operationalization • Judgment • Useful Value

Project: Terapias Bogotá | Salesforce Developer Edition + VS Code/Git | September 29, 2026

Key result

The prototype progressed from 93.33% complete accuracy in V1 to 96.67% in V2 and 30/30 correct cases in V3 on a controlled synthetic dataset. Salesforce demonstrated both accepted and human-corrected AI recommendations. A 10-case timing pilot observed average handling time decrease from 1.7 s to 1.0 s (41.18%).

Scope note: these are controlled-prototype results, not universal production claims.

1. Executive Summary

The project converts unstructured website requests into structured CRM categories while retaining final decision authority with a human reviewer. Evidence was deliberately built as an auditable chain: predefined Ground Truth, versioned prompts, measured outputs, error analysis, Salesforce implementation, governance controls, and value measurement.

Website request → AI triage → Type / Service / Confidence / Reason → Salesforce Lead → Human Review → Accept or Correct

Criterion	Evidence	Status
Interaction	V1 → V2 → V3	Complete
Evaluation	30-case Ground Truth + results	Complete
Refinement	Error + regression analyses	Complete
Operationalization	Salesforce + human review	Complete
Judgment	Risk + policy + decision log	Complete
Useful Value	Quality + timing pilot	Complete

2. Business Problem and Scope

Administrative website requests must be categorized consistently before CRM use. AI was evaluated as a recommendation layer to reduce classification effort while keeping human validation.

Allowed: Request Type, Requested Service, short administrative rationale, and confidence.
Excluded: diagnosis, treatment recommendations, clinical assessment, and medical urgency.

3. Salesforce Implementation

The standard Lead object was extended with seven custom fields and retrieved into Salesforce DX metadata for independent inspection and version control.

Field	API name	Purpose
Requested Service	Requested_Service__c	Service
Request Type	Request_Type__c	Final classification
AI Category	AI_Category__c	AI recommendation

AI Confidence	AI_Confidence__c	Confidence
AI Reason	AI_Reason__c	Rationale
AI Reviewed	AI_Reviewed__c	Human reviewed
AI Accepted	AI_Accepted__c	Human accepted

Request Type: Appointment, Information, Reschedule, Cancellation, Other. Service: Fisioterapia, Terapia a domicilio, Rehabilitación, Información general, Otro.

4. Evaluation Methodology

Thirty synthetic administrative requests were used. Expected labels were defined and versioned before V1. The same test set and Ground Truth were then reused for V1, V2, and V3, enabling direct comparison and regression detection.

cases.csv → prompt → predictions
ground_truth.csv → scoring → PASS / FAIL

Artifact	Role
cases.csv	Test messages
ground_truth.csv	Predefined expected labels
results_v1/v2/v3.csv	Version predictions
error_analysis_v1/v2.md	Failure and regression analysis
evaluation_summary.md	Final comparison

5. Interaction and Refinement

V1 achieved 28/30 complete classifications (93.33%). TB019 and TB026 exposed ambiguity between Information and Appointment for availability requests. V2 added an availability rule and improved to 29/30 (96.67%), but introduced a regression in TB007 by treating a general hours question as booking intent. V3 distinguished general schedules from availability to receive a service during a specific period and achieved 30/30 on the controlled set.

V1 93.33% → analyze TB019/TB026 → V2 96.67% → detect TB007 regression → V3 30/30 controlled set

6. Evaluation Results

Metric	V1	V2	V3
Request Type	93.33%	96.67%	100% (30/30)
Requested Service	100%	100%	100% (30/30)
Complete	93.33%	96.67%	100% (30/30)

Case	Ground Truth	V1	V2	V3
TB007	Information	Information ✓	Appointment ✗	Information ✓
TB019	Appointment	Information ✗	Appointment ✓	Appointment ✓
TB026	Appointment	Information ✗	Appointment ✓	Appointment ✓

V3 30/30 is reported only for the controlled synthetic set. It is not a universal accuracy claim.

7. Governance and Judgment

Governance artifacts: AI_Risk_Assessment.md, Human_Review_Policy.md, and Decision_Log.md.

- AI limited to administrative classification.

- Ground Truth defined before testing and never changed to improve scores.
- Full regression test after each refinement.
- AI Reviewed and AI Accepted stored separately.
- 30/30 not generalized beyond the test set.
- Prompt refinement stopped at V3 to reduce overfitting risk.

Confidence supports review but does not prove correctness. Clinical judgment remains outside the workflow.

8. Operationalization in Salesforce

Scenario	AI Category	Final Type	Confidence	Reviewed	Accepted
V3 accepted	Appointment	Appointment	0.96	TRUE	TRUE
V1 corrected	Information	Appointment	0.82	TRUE	FALSE

AI recommendation → Salesforce Lead → Human Review → ACCEPT or CORRECT

The corrected scenario uses a genuine historical V1 failure. Request_Type__c remains the final human-approved administrative classification.

9. Useful Value Evidence

Metric	Manual	AI-assisted
Cases	10	10
Total time	17 s	10 s
Average	1.7 s	1.0 s

Observed average decrease: 0.7 s per case, equivalent to 41.18% in this small timing pilot.

Quality measure	Result
Manual Request Type accuracy	8/10 = 80%
Manual Requested Service accuracy	7/10 = 70%
Complete manual accuracy	7/10 = 70%
V3 recommendations accepted during timing pilot	10/10 = 100%

Manual complete errors occurred in TB004, TB007, and TB008. They were not corrected retrospectively.

10. Limitations

- 30 synthetic evaluation cases; 10-case timing pilot.
- Single project author/reviewer.
- Reviewer had prior familiarity with the test set.
- Measured times were extremely short and may not reflect production workload.
- Developer Org and test Leads represent a controlled prototype.
- Production claims require new data, independent reviewers, privacy controls, monitoring, and approval.

11. Evidence Inventory

TestPage/ai/prompts/{trriage_v1,trriage_v2,trriage_v3}.md
 TestPage/evaluation/{cases,ground_truth,results_v1,results_v2,results_v3,value_measurement}.csv
 TestPage/evaluation/{error_analysis_v1,error_analysis_v2,evaluation_summary}.md
 TestPage/governance/{AI_Risk_Assessment,Human_Review_Policy,Decision_Log}.md
 salesforce/force-app/main/default/objects/Lead/fields/ [7 metadata files]

Git history preserves the sequence from Ground Truth through prompt versions, evaluations, refinements, governance, and measured value. Local authentication/tooling files are excluded by .gitignore.

12. Final Assessment Mapping

Criterion	Independent evidence
Interaction	Versioned V1/V2/V3 prompts
Evaluation	Predefined Ground Truth + version results
Refinement	TB019/TB026 failures + TB007 regression + rule changes
Operationalization	Stored Salesforce recommendations + human decisions
Judgment	Risk Assessment + Review Policy + Decision Log
Useful Value	93.33→96.67→30/30; 1.7→1.0 s pilot

13. Conclusion and Next Steps

The evidence demonstrates a complete controlled-prototype lifecycle: define, interact, evaluate, refine, govern, operationalize, and measure. Value came from the full workflow, not from prompting alone.

Recommended next steps: securely connect the Terapias Bogotá form to Salesforce, test new held-out cases with independent reviewers, expand validation data, build Salesforce acceptance/correction reports, and define production monitoring and rollback controls.

Artifact status: FINAL for controlled-prototype evaluation. Web-to-Lead and larger-scale testing are future improvements.